

Enable Triton on Intel® GPUs



Notices and Disclaimers

Performance varies by use, configuration and other factors. Learn more at www.Intel.com/PerformanceIndex.

Performance results are based on testing as of dates shown in configurations and may not reflect all publicly available updates. See backup for configuration details. No product or component can be absolutely secure.

Your costs and results may vary.

Intel technologies may require enabled hardware, software or service activation.

Results have been estimated or simulated.

Intel does not control or audit third-party data. You should consult other sources to evaluate accuracy.

All product plans and roadmaps are subject to change without notice.

Statements in this document that refer to future plans or expectations are forward-looking statements. These statements are based on current expectations and involve many risks and uncertainties that could cause actual results to differ materially from those expressed or implied in such statements. For more information on the factors that could cause actual results to differ materially, see our most recent earnings release and SEC filings at www.intc.com.

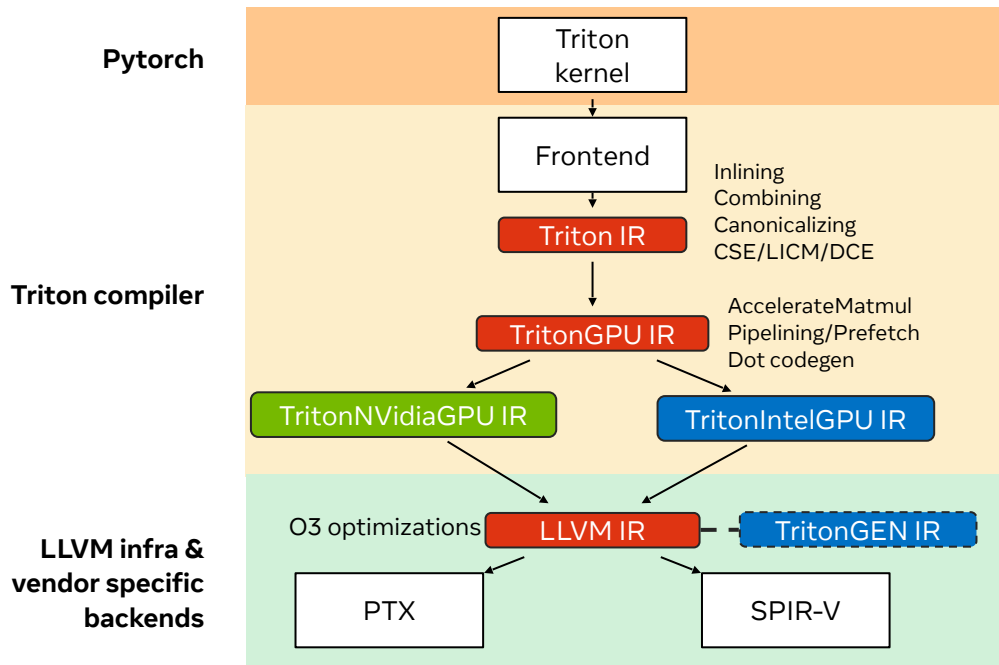
Code names are used by Intel to identify products, technologies, or services that are in development and not publicly available. These are not "commercial" names and not intended to function as trademarks.

© Intel Corporation. Intel, the Intel logo, and other Intel marks are trademarks of Intel Corporation or its subsidiaries. Other names and brands may be claimed as the property of others.

Agenda

- Intel Backend Architecture
- Functional Results & Integration with PyTorch
- HW Feature Exploitation
- Performance Results
- Challenges


Intel Backend Architecture




- Intel Triton backend redesigned (switched to lowering to the MLIR's LLVM dialect -> LLVMIR -> SPIRV)
- Reusing Triton IR passes unchanged, customized TritonGPU IR passes to support Intel's specific "Dot" operation layout (DPAS Layout)
- Intel GPUs HW exploitation abstracted by the TritonGEN dialect (e.g. DPAS, 2D block reads/writes, etc...), preempting the need to generate 'inline asm'

Functional Results & Integration with PyTorch

Intel® Data Center GPU Max



Intel® Lunar Lake GPU



- Adapted OpenAI's regression tests, approaching 99% success rate on GPU Max (aka Ponte Vecchio) and Lunar Lake (using upstream PyTorch)
- [PyTorch 2.4 added support for Intel® GPU Acceleration of AI Workloads](#) - Intel's Triton GPU backend integrated in Inductor (torch.compile mode)
- Regularly run PyTorch benchmark suites (Torchbench, Huggingface, TIMM models). On par with upstream PyTorch.

HW Feature Exploitation

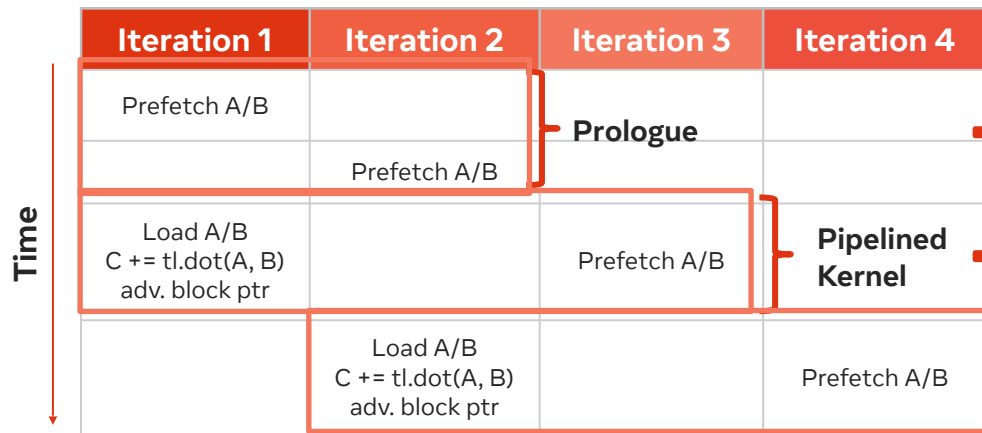


XMX Support

Data Type	Ops/Clock
TF32	2048
FP16 / BF16	4096
INT8	8192

- Intel 2nd generation Xe-core contains 8 XMX Engines supporting high performance matrix multiplications
- XMX Engine programmable via the DPAS (Dot Product Accumulate Systolic) HW instruction.
- Xe-core offers efficient HW instructions to load (and prefetch) 2D blocks from memory.

Reusing Software Pipeline Pass for Prefetch



```
%18 = tt.make_tensor_ptr %arg0, [%15, %16], [%17, %c1_i64], [%14, %c0_i32]
%22 = tt.make_tensor_ptr %arg1, [%16, %20], [%21, %c1_i64], [%c0_i32, %19]
triton_intel_gpu.prefetch %18
triton_intel_gpu.prefetch %22
%23 = tt.advance %18, [%c0_i32, %c32_i32]
%24 = tt.advance %22, [%c32_i32, %c0_i32]
triton_intel_gpu.prefetch %23
triton_intel_gpu.prefetch %24

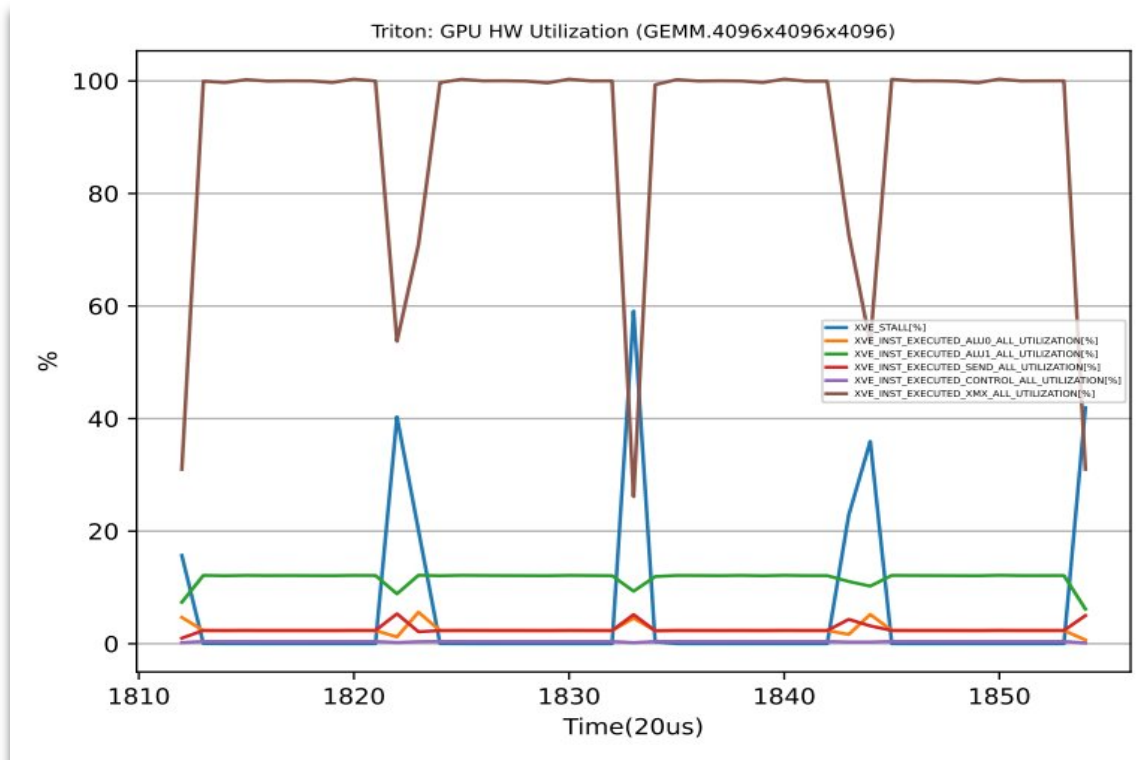
%25:5 = scf.for %arg9 = %c0_i32 to %arg5 step %c32_i32 iter_args(%arg10 = %cst, ...)
%29 = tt.advance %arg11, [%c0_i32, %c32_i32]
%30 = tt.advance %arg12, [%c32_i32, %c0_i32]
triton_intel_gpu.prefetch %29
triton_intel_gpu.prefetch %30
%31 = tt.load %arg13
%32 = tt.load %arg14
%33 = tt.dot %31, %32, %arg10
scf.yield %33, %29, %30, %arg11, %arg12
}
```

- Set up stages: prefetch, load/dot/block_ptr update
- Reuse triton passes to get prologue and pipelined kernel
- Customization: added a prefetch operation in the TritonIntelGPU dialect to prefetch to cache instead of shared memory
- Prefetch distance is tunable, trade-off between cache usage vs. latency hiding

HW Mapping


- triton_intel_gpu.prefetch [?] 2D Block Prefetches
- tt.load [?] 2D Block Loads
- tt.dot [?] DPAS (Dot Product Accumulate Systolic)

HW Utilization (GEMM)



- HW Utilization on Intel Data Center Max GPU 1550
- Achieved high XMX HW peak utilization (brown line)
- Similar utilization graph as tuned library implementation

Performance Results



- Tested on Intel Data Center Max GPU 1550
- Performance on a GEMM kernel is 90+% relative to library implementation

(*) Tested by Intel as of Aug 5, 2024. Result may vary.

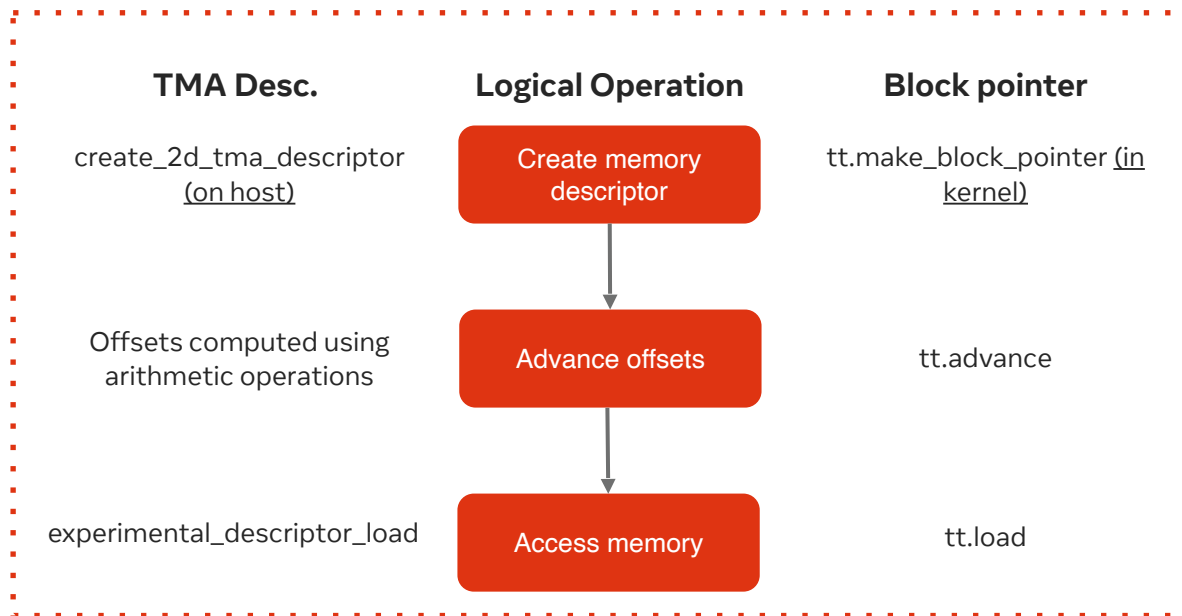
Performance Results

- Intel Triton Backend uses the semantic information associated with block pointers to exploit the 2D block reads/prefetch HW instructions
- Block pointers are a Triton experimental feature and may be deprecated upstream
- Can TMA descriptors be used in place of block pointers (to allow codegen of 2D block operations)?

TMA Descriptors vs Block pointers

- Similar layout, but TMA desc. doesn't have "block offsets"
- TMA descriptor created on the host and passed to the kernel
- Usage comparison:

TMA Descriptor	Block Pointer
Tensor Base Tensor Shapes Tensor Strides Block Shapes	Tensor Base Tensor Shapes Tensor Strides Block Shapes Block Offsets



TMA Desc. vs Block pointer

TritonIntelGPUToMLIRPass

```
tt.func public @tma_copy_kernel(%out: !tt.ptr<f32> {tt.divisibility = 16 : i32}, %tma_desc: !tt.ptr<i8, 0> {tt.nv_tma_desc = 1 : i32}) {  
  %off = arith.constant 0 : i32  
  %0 = tt.make_range [end = 128 : i32, start = 0 : i32] : tensor<128xi32, #blocked>  
  %load = tt.experimental_descriptor_load %tma_desc[%off] : !tt.ptr<i8, 0> → tensor<128xf32, #blocked>  
  %splat = tt.splat %out : !tt.ptr<f32> → tensor<128x!tt.ptr<f32>, #blocked>  
  %res = tt.addptr %splat, %0 : tensor<128x!tt.ptr<f32>, #blocked>, tensor<128xi32, #blocked>  
  tt.store %res, %load : tensor<128x!tt.ptr<f32>, #blocked>  
  tt.return  
}
```

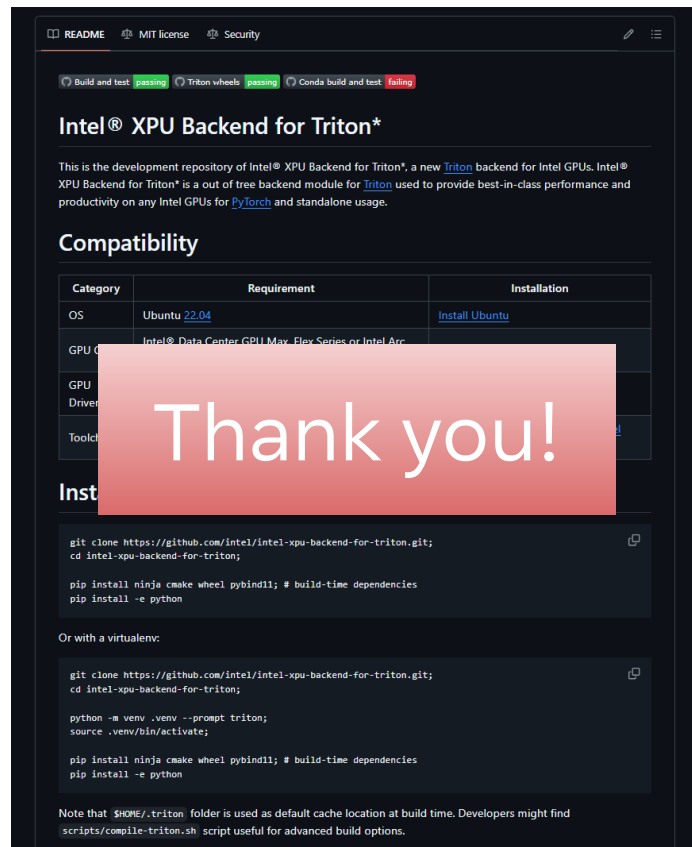
```
tt.func public @tma_copy_kernel(%out: !tt.ptr<f32> {tt.divisibility = 16 : i32}, %tma_desc: !tt.ptr<i8, 0> {tt.nv_tma_desc = 1 : i32}) {  
  %stride_one = arith.constant 1 : i64  
  %off = arith.constant 0 : i32  
  %0 = tt.make_range [end = 128 : i32, start = 0 : i32] : tensor<128xi32, #blocked>  
  %tma_desc_cast = builtin.unrealized_conversion_cast %tma_desc : !tt.ptr<i8, 0> to !llvm.ptr  
  %ptr_base = llvm.getelementptr %tma_desc_cast[0, 0] : (!llvm.ptr) → !llvm.ptr, !llvm.struct<[ptr<1>, i64, i64, i64, i64, ...]>  
  %ptr_shape = llvm.getelementptr %tma_desc_cast[0, 4] : (!llvm.ptr) → !llvm.ptr, !llvm.struct<(ptr<1>, i64, i64, i64, i64, ...)>  
  %base = llvm.load %ptr_base : !llvm.ptr → !llvm.ptr<i>  
  %base_cast = builtin.unrealized_conversion_cast %base : !llvm.ptr<i> to !tt.ptr<f32>  
  %shape = llvm.load %ptr_shape : !llvm.ptr → i64  
  %block_ptr = tt.make_tensor_ptr %base_cast, [%shape], [%stride_one], [%off] { ... } : <tensor<128xf32, #blocked>  
  %load = tt.load %block_ptr : !tt.ptr<tensor<128xf32, #blocked>  
  %splat = tt.splat %out : !tt.ptr<f32> → tensor<128x!tt.ptr<f32>, #blocked>  
  %res = tt.addptr %splat, %0 : tensor<128x!tt.ptr<f32>, #blocked>, tensor<128xi32, #blocked>  
  tt.store %res, %load : tensor<128x!tt.ptr<f32>, #blocked>  
  tt.return  
}
```

WIP

- experimental_descriptor_load “converted” to make_tensor_ptr + tt.load
- base and shape arguments obtained from the TMA descriptor
- block offset obtained from the experimental_descriptor_load operation

Project on GitHub

Triton backend for Intel® GPUs



The screenshot shows the GitHub repository page for 'Intel XPU Backend for Triton'. At the top, there are status indicators for 'Build and test' (passing), 'Triton wheels' (passing), and 'Conda build and test' (failing). The repository title is 'Intel® XPU Backend for Triton*'. Below the title, a paragraph describes it as a new Triton backend for Intel GPUs. A 'Compatibility' table is visible, listing requirements for OS, GPU, and tools. A large pink 'Thank you!' box is overlaid on the right side of the page. Below the table, the 'Installation' section provides two methods: cloning the repository and installing dependencies, and using a virtual environment.

Category	Requirement	Installation
OS	Ubuntu 22.04	Install Ubuntu
GPU	Intel® Data Center GPU Max, Flex Series or Intel Arc	
GPU Driver		
Tools		

Installation

```
git clone https://github.com/intel/intel-xpu-backend-for-triton.git;
cd intel-xpu-backend-for-triton;

pip install ninja cmake wheel pybind11; # build-time dependencies
pip install -e python
```

Or with a virtualenv:

```
git clone https://github.com/intel/intel-xpu-backend-for-triton.git;
cd intel-xpu-backend-for-triton;

python -m venv .venv --prompt triton;
source .venv/bin/activate;

pip install ninja cmake wheel pybind11; # build-time dependencies
pip install -e python
```

Note that `$HOME/.triton` folder is used as default cache location at build time. Developers might find `scripts/compile-triton.sh` script useful for advanced build options.

Configuration details

- GPU: Intel Data Center Max GPU 1550
- OS: Ubuntu [22.04](#)
- Intel Extension for PyTorch: [intel/intel-extension-for-pytorch at dev/triton-test-3.0](#)
- Intel XPU Backend for Triton: [intel/intel-xpu-backend-for-triton at a71e7be](#)
- Toolchain: [PyTorch Development Bundle v0.5.2](#)
- Graph driver: [Rolling Stable Version 914.32](#)
- Command line

```
TRITON_INTEL_ENABLE_FAST_PREFETCH=1 TRITON_INTEL_ENABLE_ADDRESS_PAYLOAD_OPT=1 IGC_VISAOptions=" -TotalGRFNum  
256 -enableBCR -nolocalra -printregusage -DPASSTokenReduction -enableHalfLSC -abiver 2" IGC_DisableLoopUnroll=1  
SYCL_PROGRAM_COMPILE_OPTIONS=" -vc-codegen -vc-disable-indvars-opt -doubleGRF -Xfinalizer ' -printregusage  
-enableBCR -DPASSTokenReduction ' "  
python benchmarks/triton_kernels_benchmark/gemm_benchmark.py
```

Creating block pointer from TMA desc.

TritonIntelGPUToMLIRLoweringPass

```
%11 = scf.for %arg3 = %c0_i32 to %c1024_i32 step %c16_i32 iter_args(%arg4 = %cst) → (tensor<128×256xf32, #blocked>) : i32 {
  %13 = tt.experimental_descriptor_load %arg0[%9, %arg3] : !tt.ptr<i8, 0> → tensor<128×16xf16, #blocked1>
  %14 = tt.experimental_descriptor_load %arg1[%arg3, %10] : !tt.ptr<i8, 0> → tensor<16×256xf16, #blocked>
  %15 = triton_gpu.convert_layout %13 : tensor<128×16xf16, #blocked1> → tensor<128×16xf16, #triton_gpu.dot_op{opIdx = 0, ...}>>
  %16 = triton_gpu.convert_layout %14 : tensor<16×256xf16, #blocked> → tensor<16×256xf16, #triton_gpu.dot_op{opIdx = 1, ...}>>
  %17 = triton_gpu.convert_layout %arg4 : tensor<128×256xf32, #blocked> → tensor<128×256xf32, #blocked2>
  %18 = tt.dot %15, %16, %17, inputPrecision = tf32
  | : tensor<128×16xf16, #triton_gpu.dot_op{opIdx = 0, ...}>> * tensor<16×256xf16, #triton_gpu.dot_op{opIdx = 1, ...}>> → tensor<128×256xf32, #blocked2>
  %19 = triton_gpu.convert_layout %18 : tensor<128×256xf32, #blocked2> → tensor<128×256xf32, #blocked>
  scf.yield %19 : tensor<128×256xf32, #blocked>
}
%12 = arith.truncf %11 : tensor<128×256xf32, #blocked> to tensor<128×256xf16, #blocked>
tt.experimental_descriptor_store %arg2[%9, %10], %12 : !tt.ptr<i8, 0>, tensor<128×256xf16, #blocked>
```

```
%11 = scf.for %arg3 = %c0_i32 to %c1024_i32 step %c16_i32 iter_args(%arg4 = %cst) → (tensor<128×256xf32, #blocked>) : i32 {
  %24 = builtin.unrealized_conversion_cast %arg0 : !tt.ptr<i8, 0> to !llvm.ptr
  // offset calculation
  %25 = llvm.getelementptr %24[0, 0] : (!llvm.ptr) → !llvm.ptr, !llvm.struct<(ptr<1>, i64, i64, i64, i64, i64, i64, i64, i64, i64, i64)>
  %26 = llvm.getelementptr %24[0, 4] : (!llvm.ptr) → !llvm.ptr, !llvm.struct<(ptr<1>, i64, i64, i64, i64, i64, i64, i64, i64, i64, i64)>
  %27 = llvm.getelementptr %24[0, 5] : (!llvm.ptr) → !llvm.ptr, !llvm.struct<(ptr<1>, i64, i64, i64, i64, i64, i64, i64, i64, i64, i64)>
  %28 = llvm.getelementptr %24[0, 9] : (!llvm.ptr) → !llvm.ptr, !llvm.struct<(ptr<1>, i64, i64, i64, i64, i64, i64, i64, i64, i64, i64)>
  %29 = llvm.load %25 : !llvm.ptr → !llvm.ptr<1>
  %30 = builtin.unrealized_conversion_cast %29 : !llvm.ptr<1> to !tt.ptr<f16>
  %31 = llvm.load %26 : !llvm.ptr → i64
  %32 = llvm.load %27 : !llvm.ptr → i64
  %33 = llvm.load %28 : !llvm.ptr → i64
  %34 = tt.make_tensor_ptr %30, [%32, %31], [%33, %c1_i64], [%9, %arg3] {order = array<i32: 1, 0>} : <tensor<128×16xf16, #blocked1>>
  %35 = tt.load %34 : !tt.ptr<tensor<128×16xf16, #blocked1>>
  ... <similarly, offset calculation for %arg1> ...
  %46 = tt.make_tensor_ptr %42, [%44, %43], [%45, %c1_i64], [%arg3, %10] {order = array<i32: 1, 0>} : <tensor<16×256xf16, #blocked>>
  %47 = tt.load %46 : !tt.ptr<tensor<16×256xf16, #blocked>>
  %48 = triton_gpu.convert_layout %35 : tensor<128×16xf16, #blocked1> → tensor<128×16xf16, #triton_gpu.dot_op{opIdx = 0, ...}>>
  %49 = triton_gpu.convert_layout %47 : tensor<16×256xf16, #blocked> → tensor<16×256xf16, #triton_gpu.dot_op{opIdx = 1, ...}>>
  %50 = triton_gpu.convert_layout %arg4 : tensor<128×256xf32, #blocked> → tensor<128×256xf32, #blocked2>
  %51 = tt.dot %48, %49, %50, inputPrecision = tf32
  | : tensor<128×16xf16, #triton_gpu.dot_op{opIdx = 0, ...}>> * tensor<16×256xf16, #triton_gpu.dot_op{opIdx = 1, ...}>> → tensor<128×256xf32, #blocked2>
  %52 = triton_gpu.convert_layout %51 : tensor<128×256xf32, #blocked2> → tensor<128×256xf32, #blocked>
  scf.yield %52 : tensor<128×256xf32, #blocked>
}
%12 = arith.truncf %11 : tensor<128×256xf32, #blocked> to tensor<128×256xf16, #blocked>
... <offset calculation for %arg2> ...
%23 = tt.make_tensor_ptr %19, [%21, %20], [%22, %c1_i64], [%9, %10] {order = array<i32: 1, 0>} : <tensor<128×256xf16, #blocked>>
tt.store %23, %12 : !tt.ptr<tensor<128×256xf16, #blocked>>
```

Thank you!

