

# Triton CPU

May 30, 2024

# Direction

- Lowering TTIR to **vector** MLIR is an effective approach. We will use it.
- What about the approach using direct vectorization?
  - It needs the same information that we need to lower to **vector** dialect
  - E.g., **tt.load** is lowered to **load <128 x float>** or **call <1024 x float> @llvm.masked.load.v1024f32.p0**.
  - If necessary, we could add another pass for manual intrinsic-based vectorization

# Lowering to vector dialect (1/2)

- Pointers

- Two styles
  - `tt.make_tensor_ptr`
  - “Idiom”: `tt.make_range + tt.splat + tt.addptr`
- `make_tensor_ptr` is easier for CPU
- In GPUs, `add_rewrite_tensor_pointer` converts it into the idiom
- Simple cases for the idiom are implemented (e.g., using `vector.splat`)
- An idea of “reverse” such idioms into `make_tensor_ptr`? Too fragile?

# Lowering to vector dialect (2/2)

- Reductions
  - Lowers to `vector.multi_reduction`
  - Fused-softmax runs

```
%11 = vector.maskedload %10[%c0], %7, %cst_3 : ... into vector<128xf32>
%12 = vector.multi_reduction <maxnumf>, %11, %cst_1 [0] : vector<128xf32> to f32
%13 = vector.splat %12 : vector<128xf32> loc(#loc10)
%14 = arith.subf %11, %13 : vector<128xf32> loc(#loc10)
%15 = math.exp %14 : vector<128xf32> loc(#loc11)
%16 = vector.multi_reduction <add>, %15, %cst_0 [0] : vector<128xf32> to f32
```

# Study: Generated Code Quality for Vector Addition

- We evaluated the quality of the generated code (**vector** dialect + LLVM)
  - Example: vector addition with different BLOCK\_SIZE, and with or without masks
  - TritonCPU AVX code examples: [link](#)
- We compared the code from C++ code with AVX intrinsics
  - LLVM code gen: <https://godbolt.org/z/f1KE7frxe>
  - GCC code gen: <https://godbolt.org/z/zjdvEqYnr>
  - LLVM generated slightly inefficient mask register spill code. GCC spills %k register to GPR.
- Findings
  - The two approaches seem to generate similar quality code
- TODO: Evaluate the quality for the reductions

# Potential fast path optimization

- Eliding mask operations?
- For certain cases, masks are not necessary except for the last iteration
- Measured expected speedups for GEMV/GEMM? Good enough to invest?

# Supporting ARM (Neon/SVE)

- We can build and run triton-cpu on a M1 machine (only Neon)
- We want to cross compile to leverage SVE
- Code quality?
  - For vector addition, the code quality looks on par with AVX
  - Masked code seems much complex, but looks okay
  - TODO: evaluate the code quality for the reductions

# Pipeline structure

- Current Nvidia/AMD have a single big pipeline with many `populate*` style
- For CPU, we attempt to use more fine-grained pipeline style



# Plans

- Feature completeness
- Performance
- Runtime
- CI
  - Intel will support for x86
  - Meta might support for ARM(?)
- Coordination: [spreadsheet](#) ⇒ Will move to Github Issues

# Contributors

- Intel: Ilya Enkovich
- Meta: Minjang Kim, Ruiqui Gao (intern), Digant Desai, Elliot Gorokhovsky