Triton Community Meeting 4/2/2024

- Upstream unit test status on AMD backend
  - AMD backend CI: MI210 (gfx90a)
    - rocm/pytorch:rocm6.0.2_ubuntu22.04_py3.10_pytorch_2.1.2
  - test_core.py failed tests
    - test_dot3d: need refactor
    - MI300 only: test_full and test_atomic_cas
  - Todo
    - Enable tests other than test_core.py
    - Add MI300X CI
- Performance gap between upstream and AMD fork on MI300
  - f8 gemm perf (tflops): 1120 (fork) vs 971 (upstream)
  - Things to add
    - Fix shared memory allocation pass
    - Add support to larger LDS vectorsize

AMD
together we advance_