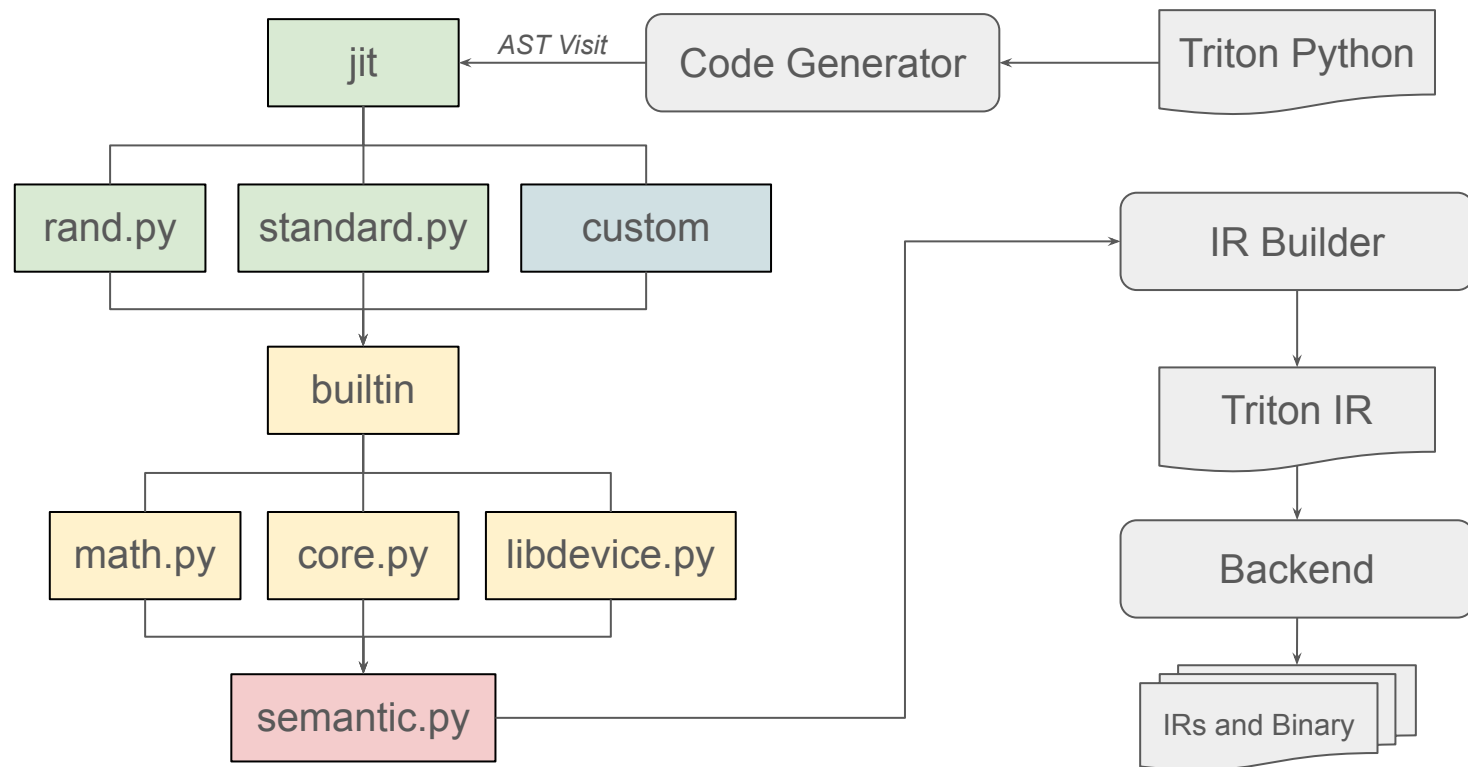


Triton Interpreter Update

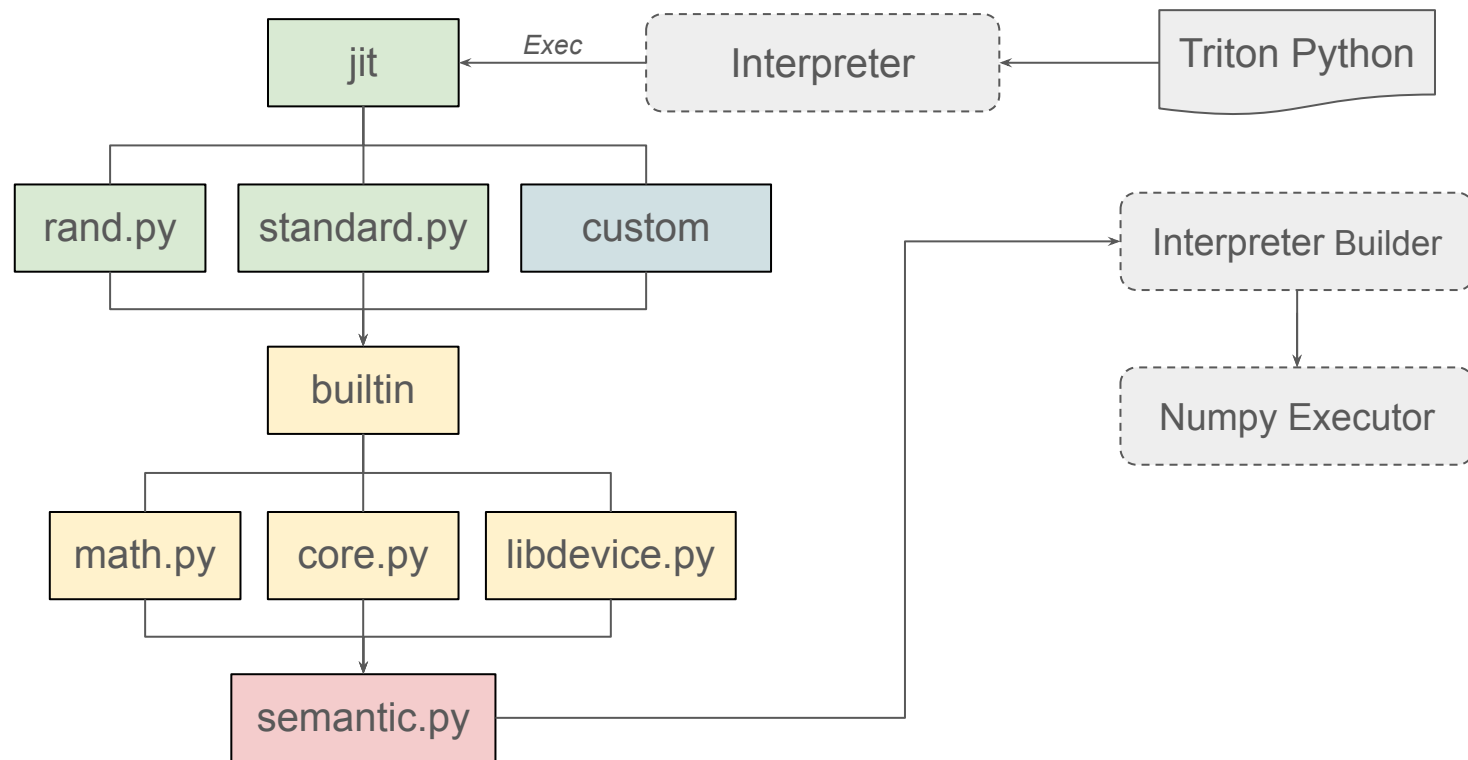
Goal

- Debug Triton code by applying print or attaching a debugger to step through the execution of individual Triton programs
 - Mostly designed for frontend users
 - Also help compiler developers get the expected output without writing a corresponding torch/pallas program
- Related files
 - *python/triton/runtime/interpreter.py*
 - *python/src/interpreter.cc*
 - *python/test/unit/language/*.py*
 - *@pytest.mark.interpreter*

Revisit the Frontend



Interpreter



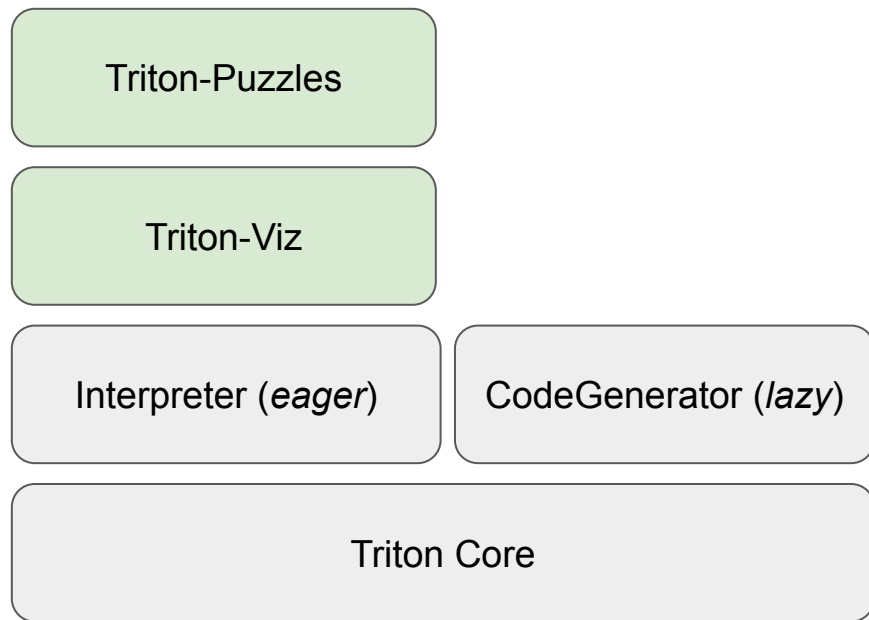
Exceptions

- `tl.reduce` and `tl.scan` are not lowered through the interpreter builder
 - `make_combine_region` invokes the code generator
 - We directly replace `tl.reduce` and `tl.scan` with custom implementations
 - Native ops like `np.sum` are accelerated through numpy
 - Custom `combine_fn` might be slow
- Functions and classes that do not go through the IR builder
 - `range`, `static_range`, `static_assert`, `static_print`
 - We replace them with python implementations

Usage


- Enable the interpreter mode
 - `TRITON_INTERPRET=1 <your command>`
- Debug with `pdb`
 - `TRITON_INTERPRET=1 pdb test.py`
 - `b test.py:<line number>`
 - `r`
- Highlights
 - You can set `device='cpu'` to execute code with the interpreter
 - You can print `tl.tensor` using the native python print and check all values of the tensor

Ecosystem



[srush/Triton-Puzzles: Puzzles for learning Triton \(github.com\)](https://github.com/srush/Triton-Puzzles)
[Deep-Learning-Profilng-Tools/triton-viz \(github.com\)](https://github.com/Deep-Learning-Profilng-Tools/triton-viz)

Coverage



Known Limitations

- No implicit **scalar** to **tensor** conversion
 - The following statements are not supported
 - `a=3`
 - `print(a.dtype)` # runtime error
 - `tl.full` # workaround1
 - `tl.to_tensor` # workaround2
- Indirect memory access is not supported
 - `ptr = tl.load(a)`
 - `a = tl.load(ptr)`
- Some precisions are not supported
 - `bfloat16`
 - `float8` series

Known Limitations

- Do not support selective interpretation
 - Only interpret all kernels
 - *Triton-Viz* doesn't have this limitation
 - Use `importlib.reload(tl)`
- Each program is executed in a fixed order
 - program id0 -> program id1 -> program id2 -> ...
- Overhead might be high
 - Especially for `tl.reduce/tl.scan` with custom associative operators
 - *Triton-Viz* mitigates the problem
 - Can sample programs

Action Items

- Documentation
- Overhead reduction
- Selective kernel interpretation?
- TorchInductor debugging?
- Case studies
 - Correctness
 - Performance estimate
- Anything else?